

CSC3022H: Machine Learning

Lab 2: K-Means Clustering

Department of Computer Science
University of Cape Town

Due: Monday, 12 August, 2019, 10.00 AM

Problem Description

Implement (in C++) the *K-means* clustering algorithm [MacQueen, 1967] with a Euclidean distance metric. See online tutorials at:

- http://www.saedsayad.com/clustering_kmeans.htm

Use the implemented K-means algorithm to cluster the following 8 examples (table 1) into 3 clusters.

When running K-means, set the initial seeds (initial centroid of each cluster) as examples 1, 4 and 7.

Table 1: Data (examples have two attributes: X , Y , both in range: $[1, 10]$).

| Example Number | X | Y |
|----------------|---|----|
| 1 | 2 | 10 |
| 2 | 2 | 5 |
| 3 | 8 | 4 |
| 4 | 5 | 8 |
| 5 | 7 | 5 |
| 6 | 6 | 4 |
| 7 | 1 | 2 |
| 8 | 4 | 9 |

Question 1: How many iterations are needed for k-means to converge?

In a text file output the results of each iteration (for each cluster, list the examples that fall into each cluster), and the centroids of each cluster, **e.g.:**

Iteration 1

Cluster 1: 1, 2, 3
Centroid: (3.0, 9.5)

Cluster 2: 4, 5, 6
Centroid: (6.5, 5.25)

Cluster 3: 7, 8
Centroid: (1.5, 3.5)

...

Iteration N

Cluster 1: 8, 7, 6
Centroid: (1.5, 3.5)

Cluster 2: 5, 4, 3
Centroid: (6.5, 5.25)

Cluster 3: 2, 1
Centroid: (3.0, 9.5)

In a ZIP file, place the source code, makefile, and the output text file (answer to question 1). Upload the ZIP file to *Vula* before 10.00 AM, Monday, 12th of August, 2019.

References

- [MacQueen, 1967] MacQueen, J. (1967). Some methods for classification and analysis of multivariate observations. In *Proceedings of the Berkeley Symposium on Mathematics, Statistics and Probability*, pages 281–297, Berkeley, USA. University of California Press.